

# EDA

```
In [1]: import numpy as np
import pandas as pd
import scipy.stats as stats
import matplotlib.pyplot as plt
import seaborn as sns

In [2]: # Loading the data from csv files.
train = pd.read_csv('train.csv')
features = pd.read_csv('features.csv')
stores = pd.read_csv('stores.csv')

In [3]: print(train.info())
print(features.info())
print(stores.info())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 421570 entries, 0 to 421569
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0    Store      421570 non-null   int64
1    Dept       421570 non-null   int64
2    Date       421570 non-null   object
3    Weekly_Sales 421570 non-null   float64
4    IsHoliday  421570 non-null   bool
dtypes: bool(1), float64(1), int64(2), object(1)
memory usage: 13.3+ MB
None
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8190 entries, 0 to 8189
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0    Store      8190 non-null   int64
1    Date       8190 non-null   object
2    Temperature 8190 non-null   float64
3    Fuel_Price 8190 non-null   float64
4    MarkDown1  4032 non-null   float64
5    MarkDown2  2921 non-null   float64
6    MarkDown3  3613 non-null   float64
7    MarkDown4  3464 non-null   float64
8    MarkDown5  4050 non-null   float64
9    CPI        7605 non-null   float64
10   Unemployment 7605 non-null   float64
11   IsHoliday  8190 non-null   bool
dtypes: bool(1), float64(9), int64(1), object(1)
memory usage: 712.0+ KB
None
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45 entries, 0 to 44
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0    Store      45 non-null   int64
1    Type       45 non-null   object
2    Size       45 non-null   int64
dtypes: int64(2), object(1)
memory usage: 1.2+ KB
None

In [4]: #Since we have been given three distinct files we joined theses and formed new data file as data.

data = train.merge(features, on=['Store', 'Date'], how='inner').merge(stores, on=['Store'], how='inner')
print(data.info())

<class 'pandas.core.frame.DataFrame'>
Int64Index: 421570 entries, 0 to 421569
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0    Store      421570 non-null   int64
1    Dept       421570 non-null   int64
2    Date       421570 non-null   object
3    Weekly_Sales 421570 non-null   float64
4    IsHoliday_x 421570 non-null   bool
5    Temperature 421570 non-null   float64
6    Fuel_Price 421570 non-null   float64
7    MarkDown1  150681 non-null   float64
8    MarkDown2  111248 non-null   float64
9    MarkDown3  137091 non-null   float64
10   MarkDown4  134967 non-null   float64
11   MarkDown5  151432 non-null   float64
12   CPI        421570 non-null   float64
13   Unemployment 421570 non-null   float64
14   IsHoliday_y 421570 non-null   bool
15   Type       421570 non-null   object
16   Size       421570 non-null   int64
dtypes: bool(2), float64(10), int64(3), object(2)
memory usage: 52.3+ MB
None

In [5]: #drop the dublicate of IsHoliday column
data = data.drop(columns=['IsHoliday_y'])
data = data.rename(columns={'IsHoliday_x': 'IsHoliday'})


In [6]: data.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 421570 entries, 0 to 421569
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  -
0    Store      421570 non-null   int64
1    Dept       421570 non-null   int64
2    Date       421570 non-null   object
3    Weekly_Sales 421570 non-null   float64
4    IsHoliday  421570 non-null   bool
5    Temperature 421570 non-null   float64
6    Fuel_Price 421570 non-null   float64
7    MarkDown1  150681 non-null   float64
8    MarkDown2  111248 non-null   float64
9    MarkDown3  137091 non-null   float64
10   MarkDown4  134967 non-null   float64
11   MarkDown5  151432 non-null   float64
12   CPI        421570 non-null   float64
13   Unemployment 421570 non-null   float64
14   Type       421570 non-null   object
15   Size       421570 non-null   int64
dtypes: bool(1), float64(10), int64(3), object(2)
memory usage: 51.9+ MB
```


In [13]: *# Pie-chart for the visual representation of store types*

```
# Data to plot
labels = 'store A','store B','store C'
sizes = [(22/(45))*100,(17/(45))*100,(6/(45))*100]
colors = ['gold', 'yellowgreen', 'lightcoral']
explode = (0.1, 0, 0) # explode 1st slice

# Plot
plt.pie(sizes, labels=labels,autopct='%1.1f%%', shadow=True, startangle=140)
plt.legend(labels, loc="best")
plt.axis('equal')
plt.show()
```




In [7]: *# boxplot for sizes of types of stores*  
store\_type = pd.concat([stores['Type'], stores['Size']], axis=1)  
f, ax = plt.subplots(figsize=(8, 6))  
fig = sns.boxplot(x='Type', y='Size', data=store\_type)



There are 45 stores in total. There are a total of 3 types of stores: Type A, B, and C. By boxplot and piechart, we can say that type A store is the largest store and C is the smallest There is no overlapped area in size among A, B, and C

In [8]: data\_ll= pd.concat([data['Dept'], data['Weekly\_Sales'], data['IsHoliday']], axis=1)  
plt.figure(figsize=(20,6))  
plt.title('Box Plot of Weekly Sales by Department and Holiday')  
fig = sns.boxplot(x='Dept', y='Weekly\_Sales', data=data\_ll, showfliers=False, hue="IsHoliday")



Sales on holiday is a little bit more than sales in not-holiday From this plot, we notice the Department with the highest sales lies between Dept 60 and 80 Total we have 421570 values for training and 115064 for testing as part of the competition. But we will work only on 421570 data as we have labels to test the performance and accuracy of models

In [20]: print(data.describe().T)

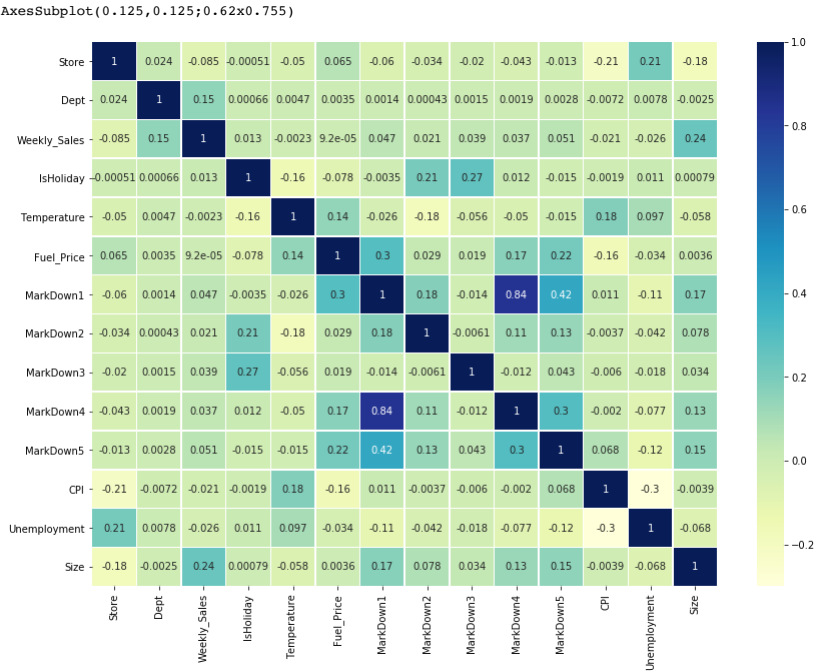
	count	mean	std	min	25%	\
Store	421570.0	22.200546	12.785297	1.000	11.000000	
Dept	421570.0	44.260317	30.492054	1.000	18.000000	
Weekly_Sales	421570.0	15981.258123	22711.183519	-4988.940	2079.650000	
Temperature	421570.0	60.090059	18.447931	-2.060	46.680000	
Fuel_Price	421570.0	3.361027	0.458515	2.472	2.933000	
MarkDown1	150681.0	7246.420196	8291.221345	0.270	2240.270000	
MarkDown2	111248.0	3334.628621	9475.357325	-265.760	41.600000	
MarkDown3	137091.0	1439.421384	9623.078290	-29.100	5.080000	
MarkDown4	134967.0	3383.168256	6292.384031	0.220	504.220000	
MarkDown5	151432.0	4628.975079	5962.887455	135.160	1878.440000	
CPI	421570.0	171.201947	39.159276	126.064	132.022667	
Unemployment	421570.0	7.960289	1.863296	3.879	6.891000	
Size	421570.0	136727.915739	60980.583328	34875.000	93638.000000	
	50%	75%	max			
Store	22.00000	33.000000	45.000000			
Dept	37.00000	74.000000	99.000000			
Weekly_Sales	7612.03000	20205.852500	693099.360000			
Temperature	62.09000	74.280000	100.140000			
Fuel_Price	3.45200	3.738000	4.468000			
MarkDown1	5347.45000	9210.900000	88646.760000			
MarkDown2	192.00000	1926.940000	104519.540000			
MarkDown3	24.60000	103.990000	141630.610000			
MarkDown4	1481.31000	3595.040000	67474.850000			
MarkDown5	3359.45000	5563.800000	108519.280000			
CPI	182.31878	212.416993	227.232807			
Unemployment	7.86600	8.572000	14.313000			
Size	140167.00000	202505.000000	219622.000000			

Replaced null values with zeros. Also,looking at the statistics of the dataframe we come to know that there are some rows for which Weekly sales have negative values. Since sales values can't be negative, we skipped those rows having negative weekly sales.

```
In [22]: data=data.fillna(0)
data = data[data['Weekly_Sales'] >= 0]
```

Out[22]: Store 0  
Dept 0  
Date 0  
Weekly\_Sales 0  
IsHoliday 0  
Temperature 0  
Fuel\_Price 0  
MarkDown1 0  
MarkDown2 0  
MarkDown3 0  
MarkDown4 0  
MarkDown5 0  
CPI 0  
Unemployment 0  
Type 0  
Size 0  
dtype: int64

```
In [35]: fig, ax = plt.subplots(figsize=(14,10)) # Sample figsize in inches
print(sns.heatmap(data.corr(),cmap="YlGnBu", annot=True, linewidths=.5, ax=ax))
```



Correlation is a bivariate analysis that measures the strength of association between two variables and the direction of the relationship. In terms of the strength of the relationship, the value of the correlation coefficient varies between +1 and -1.

```
In [38]: test = pd.read_csv('test.csv')
print(test.shape)

(115064, 4)
```

Some data preprocessing might be needed. Data is now ready for Machine learning model. However, I filled missing markdown data with zeroes. Another imputation method such as median or the mean can be tested after evaluating the model performance.

Total we have 421570 values for training and 115064 for testing as part of the competition. But we will work only on 421570 data as we have labels to test the performance and accuracy of models.